

```

# Reads the Auto data
import pandas as pd

# use pandas to read the data
auto_df = pd.read_csv("Auto.csv")

# output the first few rows
print(auto_df.head())

# output the dimensions of the data
print(auto_df.shape)

```

	mpg	cylinders	displacement	horsepower	weight	acceleration	year	\
0	18.0	8	307.0	130	3504	12.0	70.0	
1	15.0	8	350.0	165	3693	11.5	70.0	
2	18.0	8	318.0	150	3436	11.0	70.0	
3	16.0	8	304.0	150	3433	12.0	70.0	
4	17.0	8	302.0	140	3449	NaN	70.0	

	origin	name
0	1	chevrolet chevelle malibu
1	1	buick skylark 320
2	1	plymouth satellite
3	1	amc rebel sst
4	1	ford torino

```

(392, 9)

# Data exploration with code
# describe() method on the 'mpg' column with comments indicating the range and average of each column
mpg_desc = auto_df['mpg'].describe()
print("MPG column description:")
print(mpg_desc)
print("-----")
print(f"MPG column range: {mpg_desc['min']} to {mpg_desc['max']}")
print(f"Average MPG: {mpg_desc['mean']}")

# describe() method on the 'weight' column with comments indicating the range and average of each column
weight_desc = auto_df['weight'].describe()
print("\nWeight column description:")
print(weight_desc)
print("-----")
print(f"Weight column range: {weight_desc['min']} to {weight_desc['max']}")
print(f"Average weight: {weight_desc['mean']}")

# describe() method on the 'year' column with comments indicating the range and average of each column
year_desc = auto_df['year'].describe()
print("\nYear column description:")
print(year_desc)
print("-----")
print(f"Year column range: {year_desc['min']} to {year_desc['max']}")
print(f"Average year: {year_desc['mean']}")

```

```

MPG column description:
count      392.000000
mean       23.445918
std         7.805007
min         9.000000
25%        17.000000
50%        22.750000
75%        29.000000
max        46.600000
Name: mpg, dtype: float64
-----
MPG column range: 9.0 to 46.6
Average MPG: 23.445918367346938

```

```

Weight column description:
count      392.000000
mean      2977.584184
std       849.402560
min      1613.000000

```

```

25%      2225.250000
50%      2803.500000
75%      3614.750000
max       5140.000000
Name: weight, dtype: float64
-----
Weight column range: 1613.0 to 5140.0
Average weight: 2977.5841836734694

```

```

Year column description:
count      390.000000
mean       76.010256
std        3.668093
min        70.000000
25%        73.000000
50%        76.000000
75%        79.000000
max        82.000000
Name: year, dtype: float64
-----
Year column range: 70.0 to 82.0
Average year: 76.01025641025642

```

```

# Explore data types
# check the data types of all columns
print("Data types before conversion:")
print(auto_df.dtypes)

# change the cylinders column to categorical (use cat.codes)
auto_df['cylinders'] = auto_df['cylinders'].astype('category').cat.codes

# change the origin column to categorical (don't use cat.codes)
auto_df['origin'] = auto_df['origin'].astype('category')

# verify the changes with the dtypes attribute
print("\nData types after conversion:")
print(auto_df.dtypes)

```

```

Data types before conversion:
mpg          float64
cylinders      int64
displacement  float64
horsepower     int64
weight         int64
acceleration   float64
year           float64
origin         int64
name           object
dtype: object

```

```

Data types after conversion:
mpg          float64
cylinders      int8
displacement  float64
horsepower     int64
weight         int64
acceleration   float64
year           float64
origin         category
name           object
dtype: object

```

```

# Deal with NAs
# delete rows with NAs
auto_df.dropna(inplace=True)

# output the new dimensions
print("New dimensions:")
print(auto_df.shape)

```

```

New dimensions:
(389, 9)

```

```
# Modify columns
# make a new column, mpg_high, and make it categorical:
auto_df['mpg_high'] = (auto_df['mpg'] > auto_df['mpg'].mean()).astype(int)
auto_df['mpg_high'] = auto_df['mpg_high'].astype('category')

# delete the mpg and name columns (delete mpg so the algorithm doesn't just learn to predict mpg_high from mpg)
auto_df.drop(['mpg', 'name'], axis=1, inplace=True)

# output the first few rows of the modified data frame
print(auto_df.head())
```

	cylinders	displacement	horsepower	weight	acceleration	year	origin	\
0	4	307.0	130	3504	12.0	70.0	1	
1	4	350.0	165	3693	11.5	70.0	1	
2	4	318.0	150	3436	11.0	70.0	1	
3	4	304.0	150	3433	12.0	70.0	1	
6	4	454.0	220	4354	9.0	70.0	1	

	mpg_high
0	0
1	0
2	0
3	0
6	0



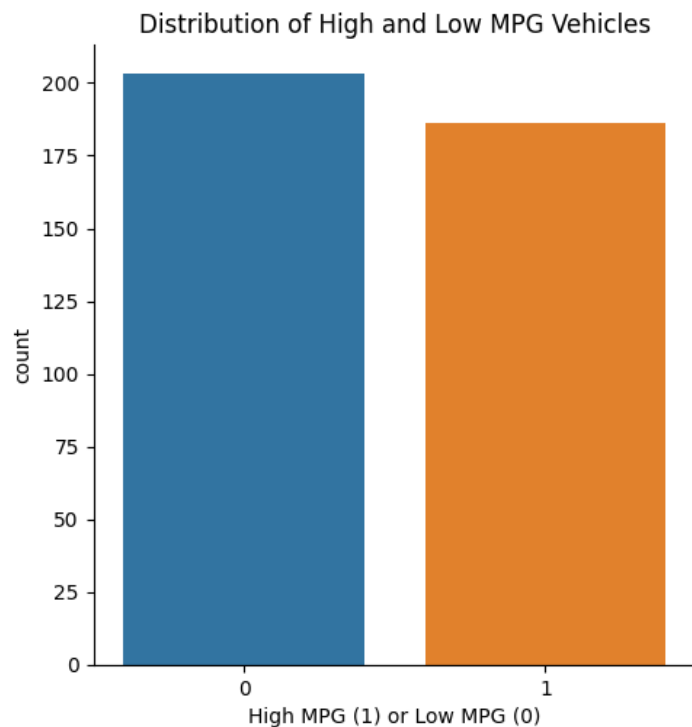
```
# Data exploration with graphs
import seaborn as sns
import matplotlib.pyplot as plt
```

```
# seaborn catplot on the mpg_high column
sns.catplot(data=auto_df, x='mpg_high', kind='count')
plt.title('Distribution of High and Low MPG Vehicles')
plt.xlabel('High MPG (1) or Low MPG (0)')
plt.show()
```

```
# seaborn relplot with horsepower on the x axis, weight on the y axis, setting hue or style to mpg_high
sns.relplot(data=auto_df, x='horsepower', y='weight', hue='mpg_high')
plt.title('Horsepower vs. Weight with High/Low MPG Highlighted')
plt.show()
```

```
# seaborn boxplot with mpg_high on the x axis and weight on the y axis
sns.boxplot(data=auto_df, x='mpg_high', y='weight')
plt.title('Distribution of Vehicle Weights by High/Low MPG')
plt.xlabel('High MPG (1) or Low MPG (0)')
plt.ylabel('Vehicle Weight')
plt.show()
```

```
# for each graph, write a comment indicating one thing you learned about the data from the graph
# Graph 1 - There are about the same amount of distribution with slightly more low than high
# Graph 2 - The distribution is linear with both weight and horse power.
# Graph 3 - In terms of distribution, the low has a higher average and range.
```



```
# Train/test split
from sklearn.model_selection import train_test_split

# 80/20
# use seed 1234 so we all get the same results
# train /test X data frames consists of all remaining columns except mpg_high
X = auto_df.drop(['mpg_high'], axis=1)
y = auto_df['mpg_high']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=1234)

# output the dimensions of train and test
```

```
print('Train dimensions:', X_train.shape, y_train.shape)
print('Test dimensions:', X_test.shape, y_test.shape)
```

```
Train dimensions: (311, 7) (311,)
Test dimensions: (78, 7) (78,)
```

```
# Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report

# train a logistic regression model using solver lbfgs
logreg = LogisticRegression(solver='lbfgs')
logreg.fit(X_train, y_train)
```

```
# test and evaluate
y_pred = logreg.predict(X_test)
```

```
# print metrics using the classification report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.80	0.88	50
1	0.73	0.96	0.83	28
accuracy			0.86	78
macro avg	0.85	0.88	0.85	78
weighted avg	0.89	0.86	0.86	78

```
/usr/local/lib/python3.9/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
# Decision Tree
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
```


```
# train a decision tree
tree = DecisionTreeClassifier(random_state=1234)
tree.fit(X_train, y_train)
```

```
# test and evaluate
y_pred = tree.predict(X_test)
```

```
# print the classification report metrics
print(classification_report(y_test, y_pred))
```

```
# plot the tree
fig, ax = plt.subplots(figsize=(12, 12))
plot_tree(tree, filled=True, feature_names=X.columns, class_names=['0', '1'], ax=ax)
plt.show()
```

	precision	recall	f1-score	support
0	0.96	0.92	0.94	50
1	0.87	0.93	0.90	28
accuracy			0.92	78
macro avg	0.91	0.92	0.92	78
weighted avg	0.93	0.92	0.92	78



```

Neural Network (15 points)
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report

train a neural network, choosing a network topology of your choice
n1 = MLPClassifier(hidden_layer_sizes=(10, 10), activation='relu', solver='adam', max_iter=1000, random_state=1234)
n1.fit(X_train, y_train)

test and evaluate
pred1 = n1.predict(X_test)

Evaluate model
print("First Neural Network:")
print(classification_report(y_test, y_pred1))

train a second network with a different topology and different settings
n2 = MLPClassifier(hidden_layer_sizes=(20, 20), activation='logistic', solver='lbfgs', max_iter=5000, random_state=1234)
n2.fit(X_train, y_train)

test and evaluate
pred2 = n2.predict(X_test)

Evaluate model

```

```
rint("Second Neural Network:")
rint(classification_report(y_test, y_pred2))
```

compare the two models and why you think the performance was same/different
there are different iterations showing and trying different methods and therefore, they differ in results.

First Neural Network:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.64	1.00	0.78	50
1	0.00	0.00	0.00	28

accuracy			0.64	78
macro avg	0.32	0.50	0.39	78
weighted avg	0.41	0.64	0.50	78

Second Neural Network:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.00	0.00	0.00	50
1	0.36	1.00	0.53	28

accuracy			0.36	78
macro avg	0.18	0.50	0.26	78
weighted avg	0.13	0.36	0.19	78

```
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.9/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision
_warn_prf(average, modifier, msg_start, len(result))
```

#	Accuracy	Recall (0)	Recall (1)	Precision (0)	Precision (1)
#Logistic Regression	0.86	0.80	0.96	0.98	0.73
#Decision Tree	0.92	0.92	0.93	0.96	0.87
#Neural Network (ReLU)	0.64	1.00	0.00	0.64	0.00
#Neural Network (logistic)	0.36	0.00	1.00	0.00	0.36

From the graph we can see that the decision tree had the best accuracy, followed by the logistic regression,
then the neural network RELU, and finally the Neural network logistic.

Fortunately for us, decision tree didn't over fit the data causing performance, therefore it was the best.
The data is linear, and therefore the logistic regression represents that.
As for Neural networks, I believe with more tuning it would perform better as they often do.

I think I might be bias on which I prefer, but I strongly prefer skit-learning. This is because at my
previous internship, I was taught to work with it. I was also taught to work with the libraries used in the project.
That being said, I do still like R, as I find it easier to do things, but skit-learning, is by far my favorite.